TREDENCE AI GRADS · PROJECT EVALUATION

Participant_01

Milestone 4 · Retail Customer Purchase
Submission:
Submission-Milestone4/Participant_M4_01_retail_purchase_strong.ipynb
Evaluated: 27 Oct 2026

89 / 100
Distinction · 89%

Criterion-wise marks

LLM Explanations & Prompt Engineering — 18 / 20

Strong prompt engineering with explicit instructions to avoid fabrication and causal claims. Could be improved by including example LLM explanation outputs demonstrating the prompt in action.

RAG: Ingestion/Retrieval/Grounding — 16 / 20

Good design and grounding principles described, including metadata preservation and retrieval-before-answering. Lacks concrete evidence of implemented ingestion, retrieval, or grounded responses with citations.

Text-to-SQL & Query Safety — 20 / 20

Excellent implementation of SQL safety checks with clear regex-based validation and comments. No evidence of schema-awareness or integration with actual retail schema, but safety enforcement is robust.

AI Agent Orchestration — 17 / 20

Strong conceptual design for agent orchestration with clear separation of concerns and traceability. Would benefit from concrete implementation or demonstration of the orchestration in action.

Evaluation & Responsible AI — 18 / 20

Comprehensive responsible AI considerations and evaluation plan are well documented. No evidence of actual testing or results, but the documented approach is thorough and appropriate.

Strengths

• Clear and explicit prompt engineering to prevent fabrication and causal claims.
• Robust SQL safety validation code enforcing read-only queries.
• Well-considered agent design separating retrieval, SQL, and summarization tools.
• Comprehensive responsible AI evaluation checklist and privacy guardrails.

To improve next time

• Provide concrete examples or outputs demonstrating LLM explanations using the prompt.
• Include actual ingestion, retrieval, and grounded response examples with citations.
• Demonstrate schema-aware SQL generation or integration with retail database schema.
• Show implemented agent orchestration workflows or code.
• Include results or logs from responsible AI tests and failure handling.

Overall feedback

The submission presents a strong conceptual and partial technical implementation for Milestone 4 requirements. It excels in prompt engineering, SQL safety, and responsible AI planning. However, it lacks concrete demonstrations of retrieval grounding, agent orchestration, and example LLM outputs. Providing these would strengthen the evidence of full milestone achievement. Human review is recommended to verify actual implementation and adherence to responsible AI principles.

AI-generated evaluation against the milestone rubric. Subject to trainer review. Each submission is evaluated once; this score is recorded.