TREDENCE AI GRADS · PROJECT EVALUATION

Participant_02

Milestone 4 · Healthcare Liver Disease
Submission:
Submission-Milestone4/Participant_M4_02_healthcare_strong.ipynb
Evaluated: 27 Oct 2026

92 / 100
Distinction · 92%

Criterion-wise marks

LLM Explanations & Prompt Engineering — 18 / 20

Strong evidence of cautious, evidence-grounded explanations and prompt engineering that aligns well with the rubric. Minor deduction as no explicit example outputs of LLM explanations are shown in the notebook to fully confirm adherence.

RAG: Ingestion/Retrieval/Grounding — 17 / 20

Good architectural design and policies for retrieval and grounding are described, but lack of concrete retrieval examples or evidence of abstention limits full marks.

Text-to-SQL & Query Safety — 20 / 20

Excellent implementation of schema-constrained, read-only SQL validation with clear rejection of unsafe commands and multiple statements.

AI Agent Orchestration — 18 / 20

Well-designed modular orchestration with clear separation of concerns and privacy considerations. Slight deduction due to absence of concrete orchestration or error handling examples.

Evaluation & Responsible AI — 19 / 20

Comprehensive coverage of responsible AI considerations and evaluation tests. Minor deduction as no actual test results or examples of these checks in action are shown.

Strengths

• Clear and cautious prompt engineering emphasizing evidence grounding and uncertainty.
• Robust SQL validation enforcing read-only queries and blocking unsafe commands.
• Thoughtful architecture separating retrieval, querying, and summarization with privacy safeguards.
• Comprehensive evaluation checklist addressing privacy, fairness, hallucination, and escalation.
• Explicit healthcare-specific guardrails to prevent misuse and clarify educational purpose.

To improve next time

• Include concrete examples of LLM explanations with source citations to demonstrate grounding.
• Provide sample retrieval outputs and abstention cases to verify RAG implementation.
• Show orchestration flows with error handling and logging in practice.
• Demonstrate evaluation checklist tests with actual results or logs.

Overall feedback

The submission presents a well-structured and safety-conscious approach to Milestone 4 requirements, with strong prompt engineering, SQL safety, and responsible AI considerations. The architectural design for retrieval and agent orchestration is sound and privacy-aware. However, the lack of concrete execution outputs, retrieval examples, and test results limits full verification of grounding, abstention, and error handling. Including such evidence would strengthen confidence in the implementation and completeness.

AI-generated evaluation against the milestone rubric. Subject to trainer review. Each submission is evaluated once; this score is recorded.