TREDENCE AI GRADS · PROJECT EVALUATION

Participant_03

Milestone 4 · Loan Default
Submission:
Submission-Milestone4/Participant_M4_03_loan_default_medium.ipynb
Evaluated: 27 Oct 2026

41 / 100
Below 50% · 41%

Criterion-wise marks

LLM Explanations & Prompt Engineering — 10 / 20

Good foundational prompt engineering and clear instructions for responsible explanations are present. To improve, demonstrate actual LLM-generated borrower-risk explanations with citations and show how unsupported claims are avoided in outputs.

RAG: Ingestion/Retrieval/Grounding — 5 / 20

The conceptual design for RAG is sound and aligns with requirements. Actual implementation and demonstration of ingestion, retrieval, and grounding with citations are needed to earn full marks.

Text-to-SQL & Query Safety — 12 / 20

Excellent safety controls for SQL queries are implemented. To improve, include examples of natural language queries converted to SQL, validation of those queries, and demonstration of safe execution.

AI Agent Orchestration — 6 / 20

The design principles for agent orchestration are appropriate and responsible. Implementation and demonstration of agent routing, error handling, and integrated workflows are needed to improve the score.

Evaluation & Responsible AI — 8 / 20

The submission shows good awareness of responsible AI principles and evaluation criteria. To improve, include executed test cases, quantitative safety evaluations, and discussion of privacy and fairness impacts.

Strengths

• Clear and responsible prompt engineering with grounding instructions.
• Strong SQL safety validation code enforcing read-only queries.
• Thoughtful design for RAG and agent orchestration emphasizing human oversight and safety.
• Explicit evaluation checklist addressing key responsible AI concerns.

To improve next time

• Demonstrate actual LLM outputs with grounded explanations and citations.
• Implement and show RAG ingestion, retrieval, and grounding with evidence.
• Provide examples of natural language to SQL translation and validation.
• Develop and demonstrate agent orchestration workflows with error handling.
• Include executed test cases and quantitative evaluations for safety, privacy, and fairness.

Overall feedback

This submission provides a well-considered conceptual framework and scaffolding for Milestone 4 requirements, including prompt design, SQL safety, RAG approach, and responsible AI evaluation. However, it lacks executed evidence such as running RAG indices, agent orchestration, actual LLM explanations, and quantitative tests. To fully meet the milestone expectations, the participant should implement and demonstrate these components with real outputs and validations. Human review is recommended due to the absence of executed results and quantitative evidence.

AI-generated evaluation against the milestone rubric. Subject to trainer review. Each submission is evaluated once; this score is recorded.