TREDENCE AI GRADS · PROJECT EVALUATION

Participant_04

Milestone 4 · Retail Customer Purchase
Submission:
Submission-Milestone4/Participant_M4_04_retail_purchase_medium.ipynb
Evaluated: 27 Oct 2026

55 / 100
Pass · 55%

Criterion-wise marks

LLM Explanations & Prompt Engineering — 12 / 20

Good prompt design and explicit instructions to avoid fabrications and causality claims. To improve, include sample LLM explanation outputs based on validated data to demonstrate the prompt's effectiveness and adherence to guidelines.

RAG: Ingestion/Retrieval/Grounding — 8 / 20

The design shows understanding of RAG principles and grounding with source metadata. To improve, provide implemented ingestion and retrieval code, demonstrate chunking, and show example grounded responses citing sources.

Text-to-SQL & Query Safety — 15 / 20

Strong safety checks for SQL queries are implemented and explained. To improve, include examples of schema-aware SQL generation, input validation, and rejection of unsafe queries in practice.

AI Agent Orchestration — 10 / 20

Conceptual design for agent orchestration is sound and aligns with best practices. To improve, provide implemented agent code, demonstrate coordination between tools, and show traceable outputs from the agent workflow.

Evaluation & Responsible AI — 10 / 20

Comprehensive evaluation and responsible AI considerations are documented. To improve, implement and report results of these tests, document limitations explicitly, and provide examples of failure handling.

Strengths

• Clear and responsible prompt engineering instructions.
• Thoughtful SQL safety validation code.
• Well-considered modular agent design and separation of concerns.
• Comprehensive evaluation checklist and privacy guardrails.

To improve next time

• Provide actual LLM explanation outputs to demonstrate prompt effectiveness.
• Implement and show ingestion, retrieval, and grounded response examples.
• Include schema-aware SQL generation examples and input validation.
• Deliver implemented agent orchestration code with traceable outputs.
• Conduct and report evaluation tests and failure cases.

Overall feedback

This submission presents a well-structured and responsible design for Milestone 4, addressing all rubric criteria conceptually with relevant code snippets and detailed plans. However, it remains largely a prototype or design document without implemented components or empirical evidence of functionality. To fully meet the milestone expectations, the participant should provide executable code for retrieval and agent orchestration, demonstrate grounded LLM explanations, and include evaluation results with documented limitations. Human review is required to verify the implementation and adherence to responsible AI principles.

AI-generated evaluation against the milestone rubric. Subject to trainer review. Each submission is evaluated once; this score is recorded.