TREDENCE AI GRADS · PROJECT EVALUATION

Participant_05

Milestone 4 · Healthcare Liver Disease
Submission:
Submission-Milestone4/Participant_M4_05_healthcare_weak.ipynb
Evaluated: 27 Oct 2026

3 / 100
Below 50% · 3%

Criterion-wise marks

LLM Explanations & Prompt Engineering — 2 / 20

The submission outlines intentions for cautious, evidence-grounded explanations but lacks concrete examples or outputs demonstrating this. To improve, include sample LLM-generated explanations with clear disclaimers distinguishing associations from diagnosis or treatment advice.

RAG: Ingestion/Retrieval/Grounding — 0 / 20

No evidence of implemented RAG pipeline or retrieval with traceable citations is present. To improve, demonstrate ingestion of approved documents, retrieval of relevant passages with source citations, and abstention when evidence is unavailable.

Text-to-SQL & Query Safety — 0 / 20

No SQL queries or evidence of schema-constrained, read-only query usage is shown. To improve, include example safe SQL queries with schema validation and demonstrate rejection of unsafe queries or minimization of sensitive data exposure.

AI Agent Orchestration — 0 / 20

No evidence of coordinating retrieval, data query, and summarization tools with clear boundaries or error handling is provided. To improve, implement and demonstrate orchestration workflows with modular components and robust error management.

Evaluation & Responsible AI — 1 / 20

While the safety plan acknowledges key responsible AI concerns, no concrete evaluation or testing evidence is presented. To improve, include tests for grounding accuracy, privacy safeguards, fairness assessments, uncertainty quantification, and explicit disclaimers requiring clinical review.

Strengths

• Clear high-level design and safety plan acknowledging milestone requirements and responsible AI principles.
• Explicit recognition of the need for evidence-based summaries, read-only SQL, and separation of tasks.

To improve next time

• Provide concrete implementations and outputs for RAG ingestion, retrieval, and grounding with source citations.
• Include example safe SQL queries demonstrating schema constraints and privacy protection.
• Demonstrate AI agent orchestration with modular tool coordination and error handling.
• Present actual LLM explanations and prompt engineering examples that distinguish associations from clinical advice.
• Conduct and report evaluations addressing grounding, privacy, fairness, uncertainty, and escalation.

Overall feedback

The submission outlines a thoughtful high-level design and safety considerations aligned with Milestone 4 requirements but lacks any implemented code, outputs, or evaluation results. It reads as a planning document rather than a completed milestone. To meet the rubric criteria fully, the participant should provide working RAG pipelines with traceable citations, safe SQL queries, orchestrated AI agent workflows, concrete LLM explanations with disclaimers, and thorough responsible AI evaluations. Without these, the submission cannot demonstrate achievement of the milestone objectives.

AI-generated evaluation against the milestone rubric. Subject to trainer review. Each submission is evaluated once; this score is recorded.